

LDA | QDA | SVM | polynomial Regression | B-spline | natural cubic spline

Yimei Fan

Create a binary variable `crim01`, and regenerate the training set and test set with ratio 2:1.

```
library(MASS)
Boston$crim01[Boston$crim >= median(Boston$crim)] = 1
Boston$crim01[Boston$crim < median(Boston$crim)] = 0

bos_new = Boston[, c("crim01", "indus", "nox", "age", "dis", "rad", "tax")]

train = sample(1: nrow(bos_new), 2/3*nrow(bos_new))
test = (-train)

train_set = bos_new[train,]
test_set = bos_new[test,]
```

Perform LDA on the training data.

The accuracy is 0.852071.

False Positive Rate: 0.2.

False Negative Rate: 0.08860759.

```
lda.fit = lda(crim01~., data = train_set)
lda.pred = predict(lda.fit, test_set)

lda.class = lda.pred$class

table(lda.class, test_set$crim01)
```

```
##
## lda.class  0  1
##           0 87 13
##           1 10 59
```

```
mean(lda.class == test_set$crim01)
```

```
## [1] 0.8639053
```

```
#False Positive Rate
18/(72+18)
```

```
## [1] 0.2
```

```
#False Negative Rate  
7/(7+72)
```

```
## [1] 0.08860759
```

Perform QDA on the training data.

According to (b), except for `tax`, all other variables are well-associated with `crim01`.

The accuracy is 0.887574.

The False Positive Rate is 0.1979167.

The False Negative Rate is 0.02739726.

```
qda.fit = qda(crim01~ indus + nox + age + dis + rad,data = train_set)  
qda.pred = predict(qda.fit,test_set)  
  
qda.class = qda.pred$class  
  
table(qda.class, test_set$crim01)
```

```
##  
## qda.class  0  1  
##           0 94 14  
##           1  3 58
```

```
mean(qda.class == test_set$crim01)
```

```
## [1] 0.8994083
```

```
#False Positive Rate  
19/(77+19)
```

```
## [1] 0.1979167
```

```
#False Negative Rate  
2/(2+71)
```

```
## [1] 0.02739726
```

Perform KNN on the training data using cross validation to decide \$k\$.

From cross-validation result, k =1 will give us the highest accuracy, which is 0.9151796.

```
train.x = train_set[,-1]  
test.x = test_set[,-1]  
train.y = train_set[,1]  
test.y = test_set[,1]  
  
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
set.seed(100)
trControl = trainControl(method = "cv", number = 5)
fit_knn_cv = train(as.factor(crim01)~indus+ nox + age + dis + rad +tax, method = "knn",
  tuneGrid = expand.grid(k = 1:10),
  trControl = trControl,
  metric = "Accuracy",
  data = bos_new)
fit_knn_cv
```

```
## k-Nearest Neighbors
```

```
##
```

```
## 506 samples
```

```
## 6 predictor
```

```
## 2 classes: '0', '1'
```

```
##
```

```
## No pre-processing
```

```
## Resampling: Cross-Validated (5 fold)
```

```
## Summary of sample sizes: 404, 405, 404, 406, 405
```

```
## Resampling results across tuning parameters:
```

```
##
```

k	Accuracy	Kappa
1	0.9151796	0.8304011
2	0.8815535	0.7631795
3	0.9091990	0.8184131
4	0.9111404	0.8222900
5	0.9091800	0.8183498
6	0.9051800	0.8103685
7	0.9111602	0.8223192
8	0.9092192	0.8184314
9	0.9051800	0.8103529
10	0.8992976	0.7985882

```
##
```

```
## Accuracy was used to select the optimal model using the largest value.
```

```
## The final value used for the model was k = 1.
```

```
knn_best = knn(train.x, test.x, train.y, k=1)
```

Fit a support vector classifier to the training data using cross validation.

cost = 5 results in lowest cross-validation error rate, which is 0.1221925. The accuracy of the best model is 0.9151796.

```
set.seed(100)
tune.out = tune(svm, as.factor(crim01)~., data = train_set, kernel = "linear", ranges = list(cost = 1:5))
summary(tune.out)
```

```
##
```

```
## Parameter tuning of 'svm':
```

```
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     50
##
## - best performance: 0.1722816
##
## - Detailed performance results:
##       cost      error dispersion
## 1    0.001 0.2016934 0.06478989
## 2    0.010 0.2045455 0.07749428
## 3    0.100 0.1986631 0.06163918
## 4    1.000 0.1723708 0.06153277
## 5    5.000 0.1812834 0.07051983
## 6   10.000 0.1812834 0.06122648
## 7   15.000 0.1812834 0.06122648
## 8   20.000 0.1812834 0.06122648
## 9   50.000 0.1722816 0.05616946
## 10 100.000 0.1723708 0.06467284
```

```
best = tune.out$best.model
pred_svm = predict(best, test_set)

table(predict = pred_svm, truth = test_set$crim01)
```

```
##       truth
## predict 0  1
##       0 88  7
##       1  9 65
```

```
(81+72)/(81+72+9+7)
```

```
## [1] 0.9053254
```

when using "polynomial" kernels, cost = 0.1 results in lowest cv error rate, which is 0.5164884, and the accuracy is 0.9053254. When using "radial" kernels, cost = 100 results in lowest cv error rate, which is 0.06247772, and the accuracy is 0.9053254.

```
set.seed(100)
tune_poly = tune(svm, as.factor(crim01) ~ ., data = train_set, kernel = "polynomial", ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 20, 50, 100)))

set.seed(100)
tune_radial = tune(svm, as.factor(crim01) ~ ., data = train_set, kernel = "radial", ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 20, 50, 100)))

tune_radial$best.parameters$degree
```

```
## NULL
```

```
par(mfrow = c(2,3))
summary(tune_poly)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   0.1
##
## - best performance: 0.4634581
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-01 0.4634581 0.1026758
## 2 1e+00 0.4634581 0.1026758
## 3 1e+01 0.4634581 0.1026758
## 4 1e+02 0.4634581 0.1026758
## 5 1e+03 0.4634581 0.1026758
```

```
summary(tune_radial)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   1
##
## - best performance: 0.07139037
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-01 0.16051693 0.05574580
## 2 1e+00 0.07139037 0.02919475
## 3 1e+01 0.08012478 0.03450182
## 4 1e+02 0.08021390 0.04507921
## 5 1e+03 0.08636364 0.05242321
```

```
table(predict(tune_poly$best.model, test_set), test_set$crim01)
```

```
##
##      0  1
## 0  0  0
## 1 97 72
```

```
#81/88
```

```
table(predict(tune_radial$best.model, test_set),test_set$crim01)
```


```
##
##      0  1
##    0 93  2
##    1  4 70
```

```
 #(84+75)/(84+75+4+6)
```

Fit the svm model with kernel = "linear".


```
svm_l = svm(as.factor(crim01)~.,data = train_set,kernel = "linear", cost = 0.01,scale = FALSE)
par(mfrow = c(2,2))
plot(svm_l, train_set,indus ~ nox)
```

SVM classification plot




```
plot(svm_l, train_set, age ~dis)
```

SVM classification plot



```
plot(svm_1, train_set, rad ~ tax)
```

SVM classification plot




Fit the svm model with kernel = "polynomial"

```
svm_p = svm(as.factor(crim01) ~ ., data = train_set, kernel = "polynomial", cost = 0.1, scale = FALSE)


par(mfrow = c(2,2))
plot(svm_p, train_set, indus ~ nox)
```


SVM classification plot



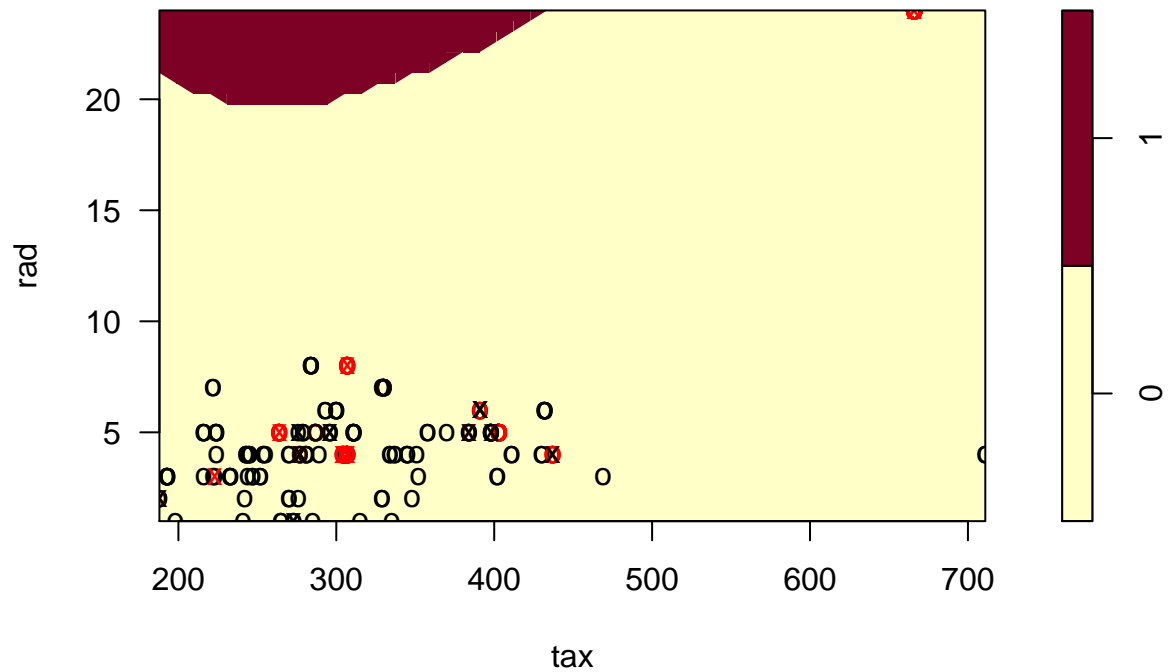
```
plot(svm_p, train_set, age ~dis)
```

SVM classification plot



```
plot(svm_p, train_set, rad ~tax)
```

SVM classification plot




Fit the svm model with kernel = "radial"

```
svm_r = svm(as.factor(crim01)~.,data = train_set,kernel = "radial", cost = 100,scale = FALSE)


par(mfrow = c(2,2))
plot(svm_r, train_set,indus ~ nox)
```

SVM classification plot



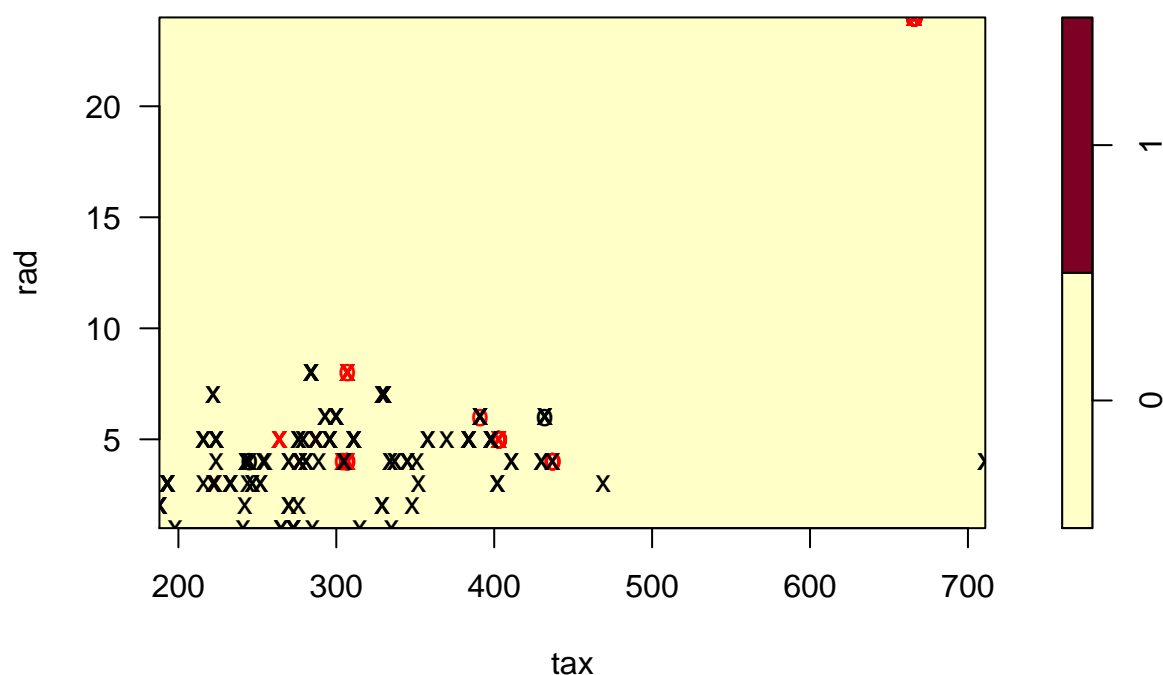
```
plot(svm_r, train_set, age ~dis)
```

SVM classification plot



```
plot(svm_r, train_set, rad ~tax)
```

SVM classification plot



(h) Compare the test errors of LDA, QDA, the best tuned models for KNN, linear SVM, SVM with radial basis

LDA: $1 - 0.852071 = 0.147929$

QDA: $1 - 0.887574 = 0.112426$

KNN: $1 - 0.9151796 = 0.0848204$

linear SVM: $1 - 0.9053254 = 0.0946746$

SVM with polynomial basis kernel : $1 - 0.9204545 = 0.0795455$

SVM with radial basis kernel : $1 - 0.9408284 = 0.0591716$

The SVM model with radial basis kernel can achieve lowest test error, while LDA can achieve highest test

Fit a cubic polynomial regression to predict `displacement` using `horsepower`.

The model is $\text{displacement} = 194.4120 + 1856.6053 \cdot \text{horsepower} - 148.3240 \cdot \text{horsepower}^2 - 224.5867 \cdot \text{horsepower}^3$

```
poly_fit = lm(displacement~poly(horsepower,3), data = Auto)
summary(poly_fit)
```


```
##
## Call:
## lm(formula = displacement ~ poly(horsepower, 3), data = Auto)
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -129.656  -21.818   -3.818   26.498  155.207
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      194.412      2.239  86.847 < 2e-16 ***
## poly(horsepower, 3)1 1856.605     44.321  41.890 < 2e-16 ***
## poly(horsepower, 3)2 -148.324     44.321  -3.347 0.000898 ***
## poly(horsepower, 3)3 -224.587     44.321  -5.067 6.26e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 44.32 on 388 degrees of freedom
## Multiple R-squared:  0.822, Adjusted R-squared:  0.8206
## F-statistic: 597.2 on 3 and 388 DF, p-value: < 2.2e-16
```

```
horselim = range(Auto$horsepower)
horse.grid = seq(from = horselim[1], to = horselim[2])
pred_poly = predict(poly_fit, newdata = list(horsepower = horse.grid), se = TRUE)
bands =cbind(pred_poly$fit +2* pred_poly$se.fit ,pred_poly$fit -2* pred_poly$se.fit)

par(mar=c(4.5,4.5,1,1) ,oma=c(0,0,4,0))
plot(Auto$horsepower ,Auto$displacement ,xlim= range(Auto$horsepower),cex =.5,col=" darkgrey ")
title(" Degree -3 Polynomial ",outer=T)
lines(horse.grid ,pred_poly$fit ,lwd=2,col="blue")
matlines (horse.grid ,bands ,lwd=1, col=" blue",lty=3)
```

Degree -3 Polynomial



Here are the plots for polynomial fits for a range of different polynomial degrees.

```
vec = c()

for (i in 1:10){

  poly_fit = lm(displacement~poly(horsepower,i), data = Auto)


  horselim = range(Auto$horsepower)
  horse.grid = seq(from = horselim[1], to = horselim[2])
  pred_poly = predict(poly_fit, newdata = list(horsepower = horse.grid), se = TRUE)
  bands =cbind(pred_poly$fit +2* pred_poly$se.fit ,pred_poly$fit -2* pred_poly$se.fit)

  par(mar=c(4.5,4.5,1,1) ,oma=c(0,0,4,0))
  plot(Auto$horsepower ,Auto$displacement ,xlim= range(Auto$horsepower),cex =.5,col=" darkgrey ")
  title(paste(" Degree - ",i," Polynomial "),outer=T)
  lines(horse.grid ,pred_poly$fit ,lwd=2,col="blue")
  matlines (horse.grid ,bands ,lwd=1, col=" blue",lty=3)


  vec[i] = sum((resid(poly_fit))^2)

}
```



Degree – 1 Polynomial




Degree – 2 Polynomial




Degree – 3 Polynomial




Degree – 4 Polynomial




Degree – 5 Polynomial




Degree – 6 Polynomial




Degree – 7 Polynomial




Degree – 8 Polynomial



Degree – 9 Polynomial



Degree – 10 Polynomial



```
print(vec)
```

```
## [1] 834610.5 812610.5 762171.3 751387.0 748198.0 746313.1 738708.4 736536.6
## [9] 736535.5 735121.6
```

Using cross-validation to select the optimal degree for the polynomial.

From the results, when degree for polynomial is 6, the model can achieve lowest sum of square error, wh

When knot is 9, the sum of square error is smallest, which is 53371.92.

```
ssr_bs = c()
for (i in 1:10){

  set.seed(12)
  testIdx = which(folds == i, arr.ind = TRUE)
  testData_bs = Auto[testIdx,]
  trainData_bs = Auto[-testIdx,]

  fit_bs = lm(displacement ~ bs(horsepower, df = 4, knots = c(i), degree = 3, intercept = FALSE), da
```

```

predd_bs = predict(fit_bs, newdata = testData_bs, se = T)

ssr_bs[i] = sum((predd_bs$fit-testData_bs$displacement)^2)
}

```

```

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in bs(horsepower, degree = 3L, knots = 3L, Boundary.knots = c(46, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

## Warning in predict.lm(fit_bs, newdata = testData_bs, se = T): prediction from a
## rank-deficient fit may be misleading

```

```
ssr_bs
```

```

## [1] 74774.23 84923.59 64515.42 69066.66 105084.51 56636.71 89553.52
## [8] 80334.43 48581.00 100365.81

```

```

fit_B = lm(displacement ~ bs(horsepower, df = 4, knots = 9, degree = 3, intercept = FALSE), data = A
pred_B = predict(fit_B, newdata = list(horsepower = horse.grid), se = T)

```

```


## Warning in predict.lm(fit_B, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```

```

plot(Auto$horsepower,Auto$displacement, col = "grey")
lines(horse.grid,pred_B$fit,lwd=2)
lines(horse.grid,pred_B$fit+2*pred_B$se,lty="dashed")
lines(horse.grid,pred_B$fit-2*pred_B$se,lty="dashed")

```



Fit a cubic B-Spline for a range of degrees of freedom from 1 to 10.

```

df_ssr = c()
for (i in 1:10){
  fit_B1 = lm(displacement ~ bs(horsepower, df = i, knots = 1,degree = 3, intercept = FALSE), data = Auto)
  pred_B1 = predict(fit_B1, newdata = list(horsepower = horse.grid), se = T)
  plot(Auto$horsepower,Auto$displacement, col = "grey")
  lines(horse.grid,pred_B1$fit,lwd=2)
  lines(horse.grid,pred_B1$fit+2*pred_B1$se,lty="dashed")
  lines(horse.grid,pred_B1$fit-2*pred_B1$se,lty="dashed")


  df_ssr[i] = sum(resid(fit_B1)^2)
}

```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```



```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
df_ssr
```

```
## [1] 762171.3 762171.3 762171.3 762171.3 762171.3 762171.3 762171.3 762171.3
## [9] 762171.3 762171.3
```

Perform cross-validation to select the best degrees of freedom for a cubic B-Spline.

When degree of freedom equals to 9, the sum of square error is smallest, which is 53371.92.

```
ssr_bs1 = c()
for (i in 1:10){

  set.seed(11)
  testIdx = which(folds == i, arr.ind = TRUE)
  testData_bs1 = Auto[testIdx,]
  trainData_bs1 = Auto[-testIdx,]

  fit_bs1 = lm(displacement ~ bs(horsepower, df = i, knots = 5, degree = 3, intercept = FALSE), data = trainData_bs1)

  predd_bs1 = predict(fit_bs1, newdata = testData_bs1, se = T)

  ssr_bs[i] = sum((predd_bs1$fit - testData_bs1$displacement)^2)
}
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in bs(horsepower, degree = 3L, knots = 5, Boundary.knots = c(46, : some
## 'x' values beyond boundary knots may cause ill-conditioned bases
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_bs1, newdata = testData_bs1, se = T): prediction from
## a rank-deficient fit may be misleading
```


```
ssr_bs
```

```
## [1] 74774.23 84923.59 64515.42 69066.66 105084.51 56636.71 89553.52
## [8] 80334.43 48581.00 100365.81
```

```
fit_B1 = lm(displacement ~ bs(horsepower, df = 5, knots = 5, degree = 3, intercept = FALSE), data = Auto)
pred_B1 = predict(fit_B1, newdata = list(horsepower = horse.grid), se = T)
```

```
## Warning in predict.lm(fit_B1, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
plot(Auto$horsepower, Auto$displacement, col = "grey")
lines(horse.grid, pred_B1$fit, lwd=2)
lines(horse.grid, pred_B1$fit+2*pred_B1$se, lty="dashed")
lines(horse.grid, pred_B1$fit-2*pred_B1$se, lty="dashed")
```




Use the ``ns()`` function to fit a natural cubic spline.

The sum of squared error is 834660.8.

```
fit_ns=lm(displacement~ns(horsepower,df=4,knots = 3),data = Auto)
pred_ns=predict(fit_ns,newdata=list(horsepower=horse.grid),se=T)
```

```
## Warning in predict.lm(fit_ns, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
plot(Auto$horsepower,Auto$displacement,xlim=horselim,cex=.5,col="darkgrey")
lines(horse.grid, pred_ns$fit,col="red",lwd=2)
```



```
sum((resid(fit_ns))^2)
```

```
## [1] 834610.5
```

Fit a natural cubic spline for a range of degrees of freedom from 1 to 10.

The resulting RSS are the same, 834660.8.

```
rss = c()
for (i in 1:10){
  fitNS = lm(displacement ~ ns(horsepower, df = i, knots = 1, intercept = FALSE), data = Auto)
  predNS = predict(fitNS, newdata = list(horsepower = horse.grid), se = T)
  plot(Auto$horsepower,Auto$displacement, col = "grey")
  lines(horse.grid,predNS$fit,lwd=2)
```



```

lines(horse.grid,predNS$fit+2*pred_B1$se,lty="dashed")
lines(horse.grid,predNS$fit-2*pred_B1$se,lty="dashed")

rss[i] = sum(resid(fitNS)^2)
}

```

```

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```

```


## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```

```

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```



```

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```

```

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```

```

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

```

```
## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading

## Warning in predict.lm(fitNS, newdata = list(horsepower = horse.grid), se = T):
## prediction from a rank-deficient fit may be misleading
```

```
rss
```

```
## [1] 834610.5 834610.5 834610.5 834610.5 834610.5 834610.5 834610.5 834610.5
## [9] 834610.5 834610.5
```

Perform cross-validation to select the best degrees of freedom for a natural cubic spline.

When degree of freedom equals to 5, the sum of squared error is smallest, which is 64785.82.

```
ssr_nss = c()
for (i in 1:10){

  set.seed(11)
  testIdx = which(folds == i, arr.ind = TRUE)
  testData_nss = Auto[testIdx,]
  trainData_nss = Auto[-testIdx,]

  fit_nss = lm(displacement ~ ns(horsepower, df = i, knots = 1, intercept = FALSE), data = trainData_nss)

  predd_nss = predict(fit_nss, newdata = testData_nss, se = T)

  ssr_nss[i] = sum((predd_nss$fit - testData_nss$displacement)^2)
}
```

```
## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
```

```
## a rank-deficient fit may be misleading

## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading

## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading

## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading

## Warning in predict.lm(fit_nss, newdata = testData_nss, se = T): prediction from
## a rank-deficient fit may be misleading
```

```
ssr_nss
```

```
## [1] 85704.73 85582.74 84389.01 72397.32 104210.76 62677.81 99188.55
## [8] 80063.60 59610.95 105655.39
```
